

Clustering of handwritten words based on structural features extraction

Lorenzo Cioni, Francesco Santoni

lore.cioni@gmail.com, fsanto92@hotmail.it

19 January 2015

Abstract

In this report we will discuss our developed application for the extraction of primitive features from images of handwritten words in order to generate clusters of similar elements, in our case the names of American States, extracted from forms of the 1930' census.

Contents

1	Introduction	2
2	The pre-existing project	2
2.1	Process steps	3
2.1.1	Localization of the area of interest	3
2.1.2	Row segmentation	4
2.1.3	Post-processing and clustering	4
2.2	Problems	4
3	Features extraction	5
3.1	Structural features	6
3.1.1	Whitespace	7
3.1.2	Loop	7
3.1.3	Dot	8
3.1.4	Diagonal line	8
3.1.5	Cross	9
3.1.6	Horizontal and Vertical line	9
3.2	Dimensional features	11
3.3	Features extraction example	12
4	Evaluating distances	14
4.1	Longest Common Subsequence	14
4.2	Euclidean Distance	14
5	Clustering	15
5.1	Affinity Propagation	15

6	Results	16
7	Compiling and running notes	21
8	Conclusion	22

1 Introduction

Segmentation and clustering of large amounts of data is one of the main research fields of modern artificial intelligence.

In recent years, with the expansion and diffusion of the digitalization of physical archives both in the private and public sector these branches of artificial intelligence have become increasingly prominent for economic and financial reasons since they can easily shrink the amount of man-hours needed to achieve the tasks involved.

Related to these circumstances, the application developed in our work presents a way to group together handwritten words with the same internal structure. Our intent is to make possible through the use of our application to group images, representing the same word, contained in handwritten scanned documents in sets with a non-banal precision. From these sets a human operator can thus categorize all the contained words through the apposition of a single label per set, rather than per element, with obvious gains in time and costs.

The basis of our work is the collection of U.S. Census data of the year 1930. Our goal is to divide enrollees by state: we do so through the clustering of the handwritten words contained in the "State" field of the form.

Document segmentation and extraction of the words corresponding to the states was developed previously by two of our colleagues in the course of *Technology of Databases* for documents with a similar form to the ones of our dataset, their project is thus the starting point of our work: our main contribution to the preexisting work is in the definition of a different, more relevant, valuation method through which determine the distance between different images, in hope that. The method proposed in our paper is the utilization of *structural features*, directly connected to the stroke with which the words were written. The main problem on which we have worked is thus the extraction of features from handwritten words with the goal of creating clusters of similar elements in order to facilitate the recognition by a human agent.

2 The pre-existing project

The census page that contains the data in digital form holds a wealth of information about each person surveyed: name, gender, membership status and some secondary features such as work or the breed.

All these data are housed in a moulded grid composed of numbered rows and columns. In addition to the citizens data, each archive also contains additional information related to the censor or data relating to samples of the population.

The project's objective was finding a particular grid area, corresponding to the area containing the registered State of each person, from which extract the handwritten text.

Following the words localization, the existing work proceeds by grouping visually similar elements to form a collection of homogeneous samples, in hope that each set will then

The image shows a historical census form from the 1940 U.S. Census. The form is titled 'POPULATION SCHEDULE' and 'SEX AND AGE DISTRIBUTION'. It contains a table with columns for household number, name, date of birth, sex, race, and marital status. The data is organized into rows, each representing a household. The form also includes a section for 'SUPPLEMENTARY QUESTIONS' at the bottom, which contains various checkboxes and fields for additional information.

Figure 1: An example of a table containing census data

contain cut outs of the same State.

The goal of the work is not to group all the words corresponding to the same state in a single cluster but rather to generate clusters that contain words that belong to a single state, the major requirement of the work is thus precision in the retrieval steps.

2.1 Process steps

We can summarize the project in three basic steps:

1. Localization of the text area containing the words of interest
2. Row segmentation
3. Post-processing and clustering

The last step, corresponding to the feature extraction and clustering of images, is the point of interest of our paper: while the past work was focused on the localization and retrieval of the samples, not much thought was put on the features to be extracted preferring simplicity and efficiency to effectiveness. It is this very point that gives our paper a reason d'être: starting from the preceding, already implemented, two steps, we proceed by providing a rewritten and improved clustering method through the use of new features, these steps will be discussed later in the Section 3 and 4.

2.1.1 Localization of the area of interest

In this first phase the aim is to identify the region of the grid within which the State words are located.

First of all the black border of the document is removed, an artefact resulting from the physical scanning. Then, through a vertical projection of the pixels and the creation of an histogram, the grid columns are located and, knowing the correct column's offset regarding the beginning of the document, just the one concerned is extracted.

2.1.2 Row segmentation

After the extraction of the concerned column from the document, our colleagues proceed with row segmentation.

Similar to the previous step, but using an horizontal projection this time, the rows of the grid can be determined with some accuracy . In this case, however, it's necessary to centre the word in the extracted image: this is done by correcting the height of each row by the analysis of black spikes on the created histogram.

At the end of this phase, ideally each word is contained in an individual image.

2.1.3 Post-processing and clustering

In the post-processing phase the individual images extracted are reworked in order to remove any vertical or horizontal residual lines left from an imperfect previous cut.

This operation of "deletion" consists in setting the related black pixels to the value of 255, pure white. In case the spurious rows are intersected by the word additional steps are taken to leave the intersecting pixels to the original value. We find an example of this in Figure 4.

This allows us to extract the most significant features in the next steps of the process. To extract the features each image is divided in windows with width of 1 pixel, on these windows are then applied the functions utilized for the extraction.

The preexisting project focus was the extraction of the words, placing less importance on the clustering and thus the features themselves, moreover the project philosophy is not to place excessive burden on the operating system, from these reasons the features are relatively simplistic: for each window the features considered are the height of the first and last black pixels and the number of transitions of the pixels from black to white and vice versa. The distance between images is found through the Euclidean Distance where the considered axis are the number of transitions and the stroke height calculated through a simple subtraction of the first and last black pixels height. The distance matrix is then fed to the Affinity Propagation algorithm to create the desired clusters.




Figure 2: An image extracted from document, the red box delineates the 1 pixel wide window.

As you can see in Figure 2 representing an extracted image, the word is centered and the bottom presents a white line: the bottom border row of the grid was intersecting the word and was extracted with it, in the pre-processing phase it was thus removed to not hinder the feature extraction.

2.2 Problems

The localization of words and their segmentation has inherent persistent difficulties.

A first error source is represented by the inherent document skew. The skew may be due to a not perfectly horizontal scan of the original document, its presence reduces the

efficacy of the work in searching for rows and columns of the grid, making the segmentation impossible with the given implementation of the project.

Other source of errors is the presence in some scanned images of other census forms below the current one: black columns from the lower document are seen as belonging to the upper one, making the application recognize them as the first column of the form.

This erroneous association renders the scan unusable since the targeted state area is found from the first black column through an hard-coded distance in pixels, moreover the analysis of the scan introduces in the clustering phase incorrect images that, recognised as similar, due to the particular way the code is implemented, skew the results in favour of the application increasing the supposed precision.




Figure 3: The resulting erroneous images

The identification of extraneous lines is another source of difficulties: the large number of dashed lines present makes the correct recognition extremely hard, it isn't always possible to *clean* the images and this can cause errors in the clustering phase.




Figure 4: Segmentation and post-processing issues

In other cases the word intersects directly gridlines. Because of this we may experience a loss of information about words if the line is removed using the method described above.

Regarding the previous source of errors many have been corrected in the new implementation: through a more precise and thorough identification of the first black column the erroneous segmentation resulting in samples similar to Figure 3 has been rendered not feasible; the problems stemmed from the difficult phase of cleaning the segmented images have been skirted through the removal of the steps to clean intersecting lines altogether, the presence of a black line doesn't generate problems for the new features' detection.

A problem of the old implementation is also the way the features extraction and the pre-processing are administered: while they could all be realized within the threads that extracted the words, they are instead left for later to be processed single-thread, causing unnecessary slowing downs. This is remedied in the new implementation.

3 Features extraction

In this work we want to improve upon the past project by providing an alternate method to calculate the distances, the focus isn't any more speed and simplicity but rather even at a major cost we want to find characteristics more deeply related to the written words nature in hope they may be better suited for the recognition of similar written words.

The idea is thus to find *primitive* features, resembling the possible types of strokes used to write a word, to characterize the sample from which they are extracted in order to perform clustering on their set.

In particular due to the way the code was implemented the feature extraction follows immediately the segmentation of the image and can thus be executed in direct succession within the same threads without adding great complexity to the process.

3.1 Structural features

The features are identified by the study of a sub-portion, or window, of the images correspondent to the entries of the "State" field of the census document on which the procedure is applied.

The features we search are intrinsically characteristics of the stroke, they occupy a non-minimal space, as such they are located through the study of the area, the window, with which they are implicitly associated; this poses a problem: there is the possibility that an area may be characterized by multiple features, in this case there is no clear order of which of the multiple features comes first. Due to the absence of such *native* order the string that defines a window is maintained consistent with the other strings through the convention of generating the string with the features' identifiers always taking the same order, if present.

The chosen way to resolve the issue however presents the problem of making sliding windows not directly applicable: this means that a sample must necessarily be cut in separate windows, which are allowed to overlap, with the corresponding result that due to the random cut some features like *loops* may not be recognised in an instance and recognised in another.

Each sample is associated to a characterizing string of characters made from the ordered identifiers of the windows in the sample, each window's string is itself made from the identifier strings of the features recognized in the sample. The features that convey more information about a word or are more precisely identified, for example loops and dots, are associated to longer strings so that their presence may be better recognised.

After having recognised the presence of a feature in a given window the corresponding identifier is added to the end of the string correspondent to the window; the order in which the features are searched and thus added to the queue is decided a priori and maintained consistent during the construction of the strings.

Windows Currently we have chosen to cut the samples in windows of 40 pixels each, with the exception of the last segment of the sample that, deemed irrelevant to the end of characterization due to containing almost always white space, is simply ignored to prevent eventual troubles due to the irregular number of pixels.

These windows are spaced only 7 pixels from the preceding and succeeding ones, with the intent of creating overlap and lengthening the string that characterizes the samples. (Both the dimension of the window and the spacing between them was decided through tests to obtain the optimal efficacy of the program)



Figure 5: Overlapping windows 40 pixels wide, spaced 7 pixels

While the simplest way to create the windows in the code would be simply cutting the original image in pieces through specialized functions of the *Leptonica* library, these

steps require a great amount of memory and time. We have thus preferred to increase the complexity of the functions that search for the features, to whom we pass as a variable the original images with information about the *offset* at which to start the search and the *width* of the window. The width of the window passed has actually a non banal meaning due to the fact that different features are associated with areas of the image with varying dimensions: for example it is not sensible to search for an horizontal line and a vertical line utilizing windows with the same width due to the fact that horizontal lines realistically will require windows with great width but will not have requisites on the height.

In particular, the features with less space requirements are searched in both 20 pixels halves of the windows and their strings are appended to the image string, the whole 40 pixels window is then searched for the space-hungry features, only then their identifiers are added to the general string.




Figure 6: A window (40px) and its first half(20px)

The height of the windows used is never considered as a parameter since for all the features searching functions the height is customarily the height of the sample.

3.1.1 Whitespace

The first feature that is searched in the windows is the presence of white space. This way if a window is identified as blank there won't be a need to proceed with the search of the other features, saving time.



Figure 7: A Whitespace

A window is identified as blank if the average pixel values is below a pre-set threshold.

The string associated with the Whitespace is " ". This string is not given much consideration when calculating distances since it may bring problems if images with the same word are not properly centered in the segmented samples, while differentiating words through the presence of spaces doesn't bring many fruits.

3.1.2 Loop

The feature that represents a loop requires the search of the whole 40 pixels window. The localization of a loop starts from the search of a black pixel. The idea behind is that a point on the border of a loop is such that by looking in a decided direction (in this case the y axis) we meet two pixel's value transitions first from black to white and then from white to black. Between the two transitions for a loop to be recognised there must be a minimum number of white pixel, number that is decided discretionally in the code.



Figure 8: A loop

To recognise a loop the above condition must also be verified in the other cardinal axis. To do so, starting from the center of the supposed loop (the median value of the segment described above) we check that moving either right or left we

find a transition from white to black after a suitable number of white pixels.

The implemented method has numerous problems related to the difficulty of finding the best white spaces threshold: an excessively little minimum number of white pixels makes us recognise as loops white noise that happens while scanning images, a threshold too high forces us to discard some small but real loops. The method also does not take into account the thickness of the stroke, rendering the overall recognition harder, with appropriate threshold values though we can achieve good results.

The string associated with the Loop feature is "LL".

3.1.3 Dot

To search for Dots within the segment we first proceed locating the *connected components* inside of it. The individual components are extracted and inserted into a *box* of which we know the size and the relative position to the segment.



Figure 9: A dot

At this point we search for those boxes with dimensions between two pre-set values (minimum and maximum radius), those that meet this condition are boxes that contain most likely points. Using the connected components we can thus retrieve dots in a simple way.

The string associated with the Dot feature is ".....".

It is particularly long because the dot helps distinguish the words that contain the "i" letter and a dot is easily recognised with little error.

3.1.4 Diagonal line

The feature representing a diagonal line, both upward and downward facing, is extracted through a simple exhaustive scansion of the window for lines of connected black pixels that have an incline in a range of values and are sufficiently long.

In the case of an upward facing diagonal, given a black pixel we consider it connected to another black pixel if this second one is in one of 3 different position that are illustrated in Figure 10. The lines recognized are thus all those with gradient between 30 and 60%.



Figure 10: Given the black pixel we search for other black pixels in the 3 positions, the closer red one is given priority

At the moment the function doesn't account for the width of the lines found, thus diagonal features can be recognised in a formless blob of black pixels that is sufficiently big.

A distinction in the associated string is introduced for the diagonal features that appear in the lower or upper bottom of the window, moreover the searches for the two



Figure 11: proper Diagonal feature(a) and a formless blob that introduces errors(b), in this second both upward and downward diagonal lines are recognised

kinds of lines are separate. At most in a window the function identifies a couple of upward and a couple of downward facing lines (lower and upper parts), once one has been found it stops searching for the same type. The lines that are found are further categorised on their length if medium or long.

Many different strings are associated with the Diagonal feature, they will be properly listed in the Table 1. In particular the string associated with the "long" version of a feature is made from the "medium" version of the string at which is attached an other character: this allows to distinguish features based on their length while at the same time maintaining similarity with a shorter version of themselves to account for variation in the stroke used to write the same word.

3.1.5 Cross

The function, having found at most a lower and upper case for upward and downward diagonal lines, proceeds to confront the edge points of these lines: if they possibly intersects it extracts a *crossing* feature with distinction if the crossing happens in the lower or upper part of the window, priority is given to the bottom part if for example a bottom upward diagonal intersects an upper downward diagonal.

The problems with this sub-feature are that intersections of bottom and upper lines of the same type (e.g. both upward facing) with different inclines are not recognised, in the same way there's no consideration for intersections made from lines that are not the "primary" bottom and upper diagonals: if in a single windows are present more lower upwards diagonals and one of them other than the first intersects the recognised downward diagonal, such a crossing is ignored.



Figure 12: A cross

The crossing feature may also not necessarily be a complete intersection, in fact for recognition there just needs to be a merging of a downward and upward line.

The string associated depending if the cross is found in the upper or lower half is "X" and "x" respectively. While it may seem that a crossing is distinctive of a word and thus deserving of more "characterizing force", that is the length of the associated strings, the difficulty in recognising the feature and it's uncertainty renders an excessive weight given to its presence dangerous for clustering purposes. The feature is thus maintained for future improvements but doesn't bring immediate contributions to the application.

3.1.6 Horizontal and Vertical line

Both horizontal and vertical lines' features are extracted through a simple scan of the window.

The horizontal lines requires a double-window for their implicit characteristic. The function identifies a line if it finds a connected row or column of black pixels that has sufficient length, in particular it distinguishes the lines found in normal or long through an ulterior threshold. In particular, as obvious difference from the diagonal line, no play is left to recognise inclined lines as horizontal or vertical.

Once a fitting line has been found the function keeps searching for more until the exhaustion of the allotted space. The scansion of the window is continued after having moved a certain distance from the last line found in order not to confuse a particularly thick stroke as different separate lines. The corresponding string as thus no limitation in length other than the ones posed by the maximum number of lines that can fit in a window.




Figure 13: Horizontal(a) and Vertical(b) features

There still persists the problem that the stroke width is not fully considered so a big blob of black pixels is seen as a series of horizontal and vertical lines, at the same time such an occurrence is rare so the presence of many horizontal and vertical lines ends up distinguishing the word in itself.

Recapitulation

We present a recapitulation of the various features and their correspondent strings in Table 1.

Feature	Associated string
Whitespace	" "
Loop	"LL"
Dot	"...."
Upward facing diagonal lower half, medium	"s"
Upward facing diagonal lower half, long	"bs"
Upward facing diagonal upper half, medium	"S"
Upward facing diagonal upper half, long	"BS"
Downward facing diagonal lower half, medium	"u"
Downward facing diagonal lower half, long	"vu"
Downward facing diagonal upper half, medium	"U"
Downward facing diagonal upper half, long	"VU"
Cross upper	"X"
Cross lower	"x"
Horizontal line lower half, medium	"_"
Horizontal line lower half, long	"h-"
Horizontal line upper half, medium	"H-"
Horizontal line upper half, long	"Hh-"
Vertical line lower half, medium	"i"
Vertical line lower half, long	"Ii"
Vertical line upper half, medium	"Vi"
Vertical line upper half, long	"VIi"

Table 1: Recapitulation of the strings associated to each feature

3.2 Dimensional features

While the already proposed structural features help in categorizing and recognising words to cluster, they aren't completely able to distinguish a word from similar others, we have thus searched for ways to improve the accuracy of the application: while many optimizations have been added to the defined structural features we have found that adding another layer of features of a different kind, in this case *dimensional*, helps greatly. Moreover the complexity added is extremely limited: as can be seen later on in the Table 4 the required calculation time doesn't increase in a particularly meaningful way: the majority of the process complexity isn't located in the features' extraction but rather in the clustering and construction of the LCS-distance matrix. Thus in order to improve word clustering we combine structural features with *dimensional features* for each word, after constructing the distance matrices for both we combine them through a factor determined from testing.

These dimensional features are exactly the ones previously utilized by our colleagues in the pre-existing work: height of the stroke, calculated through the difference between highest and lowest black pixel, and number of transition of the pixels' value from black to white and vice versa.

A simple explanation on how they are extracted can be found in section 2.1.3.

3.3 Features extraction example

We present now an example of features extraction from a word sample, in this case it is presented an older implementation of our application but the ideas and process scheme are maintained constant. In particular the older version was chosen to simplify the explanation since it uses smaller windows with lesser overlap so that it is visually more self-explanatory.

We initially create, as described in the section above, a sliding window that scrolls along the image by a fixed step. Then, for each section, we extract features and generate a substring.




Figure 14: Feature extraction from *Missouri* word




Figure 15: Sliding window segmentation

- (a) Some **diagonal lines** (ascending (s) and descending (u)), at the top (S) and at the bottom (s) of the image. Generated string: "sSUusSUSu".



- (b) As the previous image and two more diagonal lines representing the "i". An **horizontal line** (H). Generated string: "sSUusSUSuHsu".



- (c) Some diagonal lines and, at the end of the window, an horizontal line. Generated string: "ssSusSsH".



- (d) Two consecutive similar character represented by diagonal lines and **vertical lines** in the middle. Generated string: "sVHsV".



(e) In that window we can find a **dot** (.). Generated string: "*ssVHs.*"



(f) The same previous dot and a little **loop** (L) at the bottom. Generated string: "*ssVs.LssH*".



(g) The same loop as previous, an horizontal line a vertical line and two diagonal. Generated string: "*LVHssu*".



(h) Generated string: "*ssVHus*".



(i) The other dot here. Generated string "*ssus.H*".



(j) An i. Generated string: "*ssuu*".



(k) Only a diagonal line. Generated string: "*s*".



(l) Empty window. Generated string: " ".



(m) Empty window. Generated string: " ".



The resulting string is:

"sSUusSUSusSUusSUSuHsussSusSsHsVHsVssVHsssVs.LssHLVHssussVHusssus.Hssuu.s"

Once those strings are generated they are combined together to generate the *structure string* associated with the word sample.

4 Evaluating distances

After the extraction of good features from the words our aim is to construct a similarity matrix between different samples so that they can be used as input to the clustering algorithm. To do that we generate a measure of distance between couples of words through the use of the *Longest Common Subsequence* algorithm in which the handled strings are constructed through the appending of conventional identifiers associated with the features found in the word in a consistent order.

4.1 Longest Common Subsequence

In order to compare the generated strings one another we must define a distance on the samples. The distance used in our work is based on the *Longest Common Subsequence* (LCS) algorithm.

The LCS algorithm has the aim of extracting from a set of sequences (in this case only two) the longest common subsequence, that is a sequence that is obtainable from both the starting sequences by deleting some elements without changing the order of the remaining ones.

LCS is a particular case of the *Edit Distance* algorithm where the only allowed operations are insertion and deletion. The distance associated with LCS in our current work is the number of insertion and deletions that must be applied to obtain the longest subsequence, in accordance with the Edit distance where the distance is calculated with the number of operations needed to morph a string in the other (in Edit Distance it's possible moreover to confer customizable costs to the substitutions).

While the LCS algorithm may require high costs when applied concurrently to high numbers of sequences, in our case there exists an easy and light implementation that exploits *dynamic programming*, in this type of implementation the cost ends up being $O(n * m)$ where n and m are the length of the compared strings.

$$LCS(X_i, Y_j) = \begin{cases} 0 & \text{if } i = 0 \text{ or } j = 0 \\ LCS(X_{i-1}, Y_{j-1}) \cup x_i & \text{if } x_i = y_j \\ \text{longest}(LCS(X_i, Y_{j-1}), LCS(X_{i-1}, Y_j)) & \text{if } x_i \neq y_j \end{cases}$$

Obtained the Longest Common Subsequence between two strings the distance between them is thus:

$$x.length() + y.length() - 2 * LcsLength$$

where x and y are the strings and the $length()$ function returns the length of a string.

4.2 Euclidean Distance

The Longest Common Subsequence is only used with strings. In order to improve the correctness of the words' distance matrix we combine LCS distance with the Euclidean

Distance obtained through the use of the *dimensional features*.

The Euclidean Distance between two samples is evaluated through the formula:

$$d = \sqrt{((a_b - a_t)(b_b - b_t))^2 - (a_c - b_c)^2}$$

where:

- a_b representing the end of the stroke a
- a_t representing the begin of the stroke a
- b_b representing the end of the stroke b
- b_t representing the begin of the stroke b
- a_c representing the number of transition (normalized) of a
- b_c representing the number of transition (normalized) of b

5 Clustering

The clustering phase consists in the categorization of the various segments in homogeneous groups, so that, at the end of the process, each cluster contains only words corresponding to the same state of birth.

To do so we utilize the previously built distance matrices obtained by the processing of the structural and dimensional features: the matrices are combined in a single one through a simple weighted sum.

$$ClusteringMatrix[i, j] = StructuralMatrix[i, j] + 0.5 * DimensionalMatrix[i, j]$$

The factor of 0.5 was determined through extensive tests.

To utilize the vast majority of the clustering algorithms we need to know the number of desired final clusters. Since our objective isn't having a particular number of clusters but rather having clusters with elements with certain properties we need to use an algorithm that hasn't the requirement: we have thus used the *Affinity Propagation* algorithm.

5.1 Affinity Propagation

Affinity Propagation is a clustering algorithm that identifies a set of *exemplars* that represents the dataset¹. The input of Affinity Propagation is the pair-wise similarities between each pair of data points, $s[i, j] \forall i, j = 1, \dots, n^2$. Any type of similarities is acceptable thus Affinity Propagation is widely applicable.

Given similarity matrix $s[i, j]$, Affinity Propagation attempts to find the exemplars that maximize the net similarity, i.e. the overall sum of similarities between all exemplars and their member data points. The process of Affinity Propagation can be viewed as a message passing process with two kinds of messages exchanged among data points: *responsibility* and *availability*.

¹Brendan J. Frey, Delbert Dueck, *Clustering by Passing Messages Between Data Points*, <http://www.sciencemag.org/>, 2007

² $s[i, j]$ for each data point is called preference and impacts the number of clusters.

Responsibility, $r[i, j]$, is a message from data point i to j that reflects the accumulated evidence for how well-suited data point j is to serve as the exemplar for data point i .

Availability, $a[i, j]$, is a message from data point j to i that reflects the accumulated evidence for how appropriate it would be for data point i to choose data point j as its exemplar. All responsibilities and availabilities are set to 0 initially, and their values are iteratively updated as follows to compute convergence values:

$$\begin{aligned} r[i, j] &= (1 - \lambda)\rho[i, j] + \lambda r[i, j] \\ a[i, j] &= (1 - \lambda)\alpha[i, j] + \lambda a[i, j] \end{aligned}$$

where λ is a damping factor introduced to avoid numerical oscillations, and $\rho[i, j]$ and $\alpha[i, j]$ are, we call, *propagating responsibility* and *propagating availability*, respectively.

$\rho[i, j]$ and $\alpha[i, j]$ are computed by the following equations:

$$\begin{aligned} \rho[i, j] &= \begin{cases} s[i, j] \max_{k=j} a[i, k] + s[i, k] & i \neq j \\ s[i, j] \max_{k=j} s[i, k] & i = j \end{cases} \\ \alpha[i, j] &= \begin{cases} \min(0, r[j, j] + \sum_{k \neq i, j} \max(0, r[k, j]) & i \neq j \\ \sum_{k \neq j} \max(0, r[k, j]) & i = j \end{cases} \end{aligned}$$

That is, messages between data points are computed from the corresponding propagating messages. The exemplar of data point i is finally defined as:

$$\arg \max r[i, j] + a[i, j] : \forall j = 1, 2, \dots, n$$

As described above, the original algorithm requires $O(n^2t)$ time to update messages, where n and t are the number of data points and the number of iterations, respectively. This incurs excessive CPU time, especially when the number of data points is large³. The Figure 16 shows how affinity propagation works⁴.

The clustering process described is then applied to the array of similarity calculated previously on features extracted from each words. Once you have run the calculation of clusters, segments are organized into individual folders to provide a visual result of the proceedings just completed. Each group is identified by a segment that represents the centroid of the cluster, which is the element to which all others in the group are closer.

6 Results

Let's explore the test phase and see the results in detail. Initially we ran debug tests on our personal PCs in order to easily modify the code, these starting tests and the tests utilized to determine the right value for the many defined constants(altogether the vast majority of the tests done) are not shown here.

All the tests shown below were performed on a single more powerful machine that has enabled us to work with much larger data in less time. The machine used is composed of two Xeon processors for a total of 16 cores 2.80 Ghz and 48 Gb of RAM.

The tests were performed on a growing number of scans, each containing a maximum of 50 lines from which we extracted the words that represent the states. For each set of scans we present three possible distances calculation: only LCS, only L1 (Euclidean

³Yasuhiro Fujiwara, Go Irie, Tomoe Kitahara, *Fast Algorithm for Affinity Propagation*, 2009

⁴Brendan J. Frey, Delbert Dueck, *op. cit.*, Figure 1, p. 974

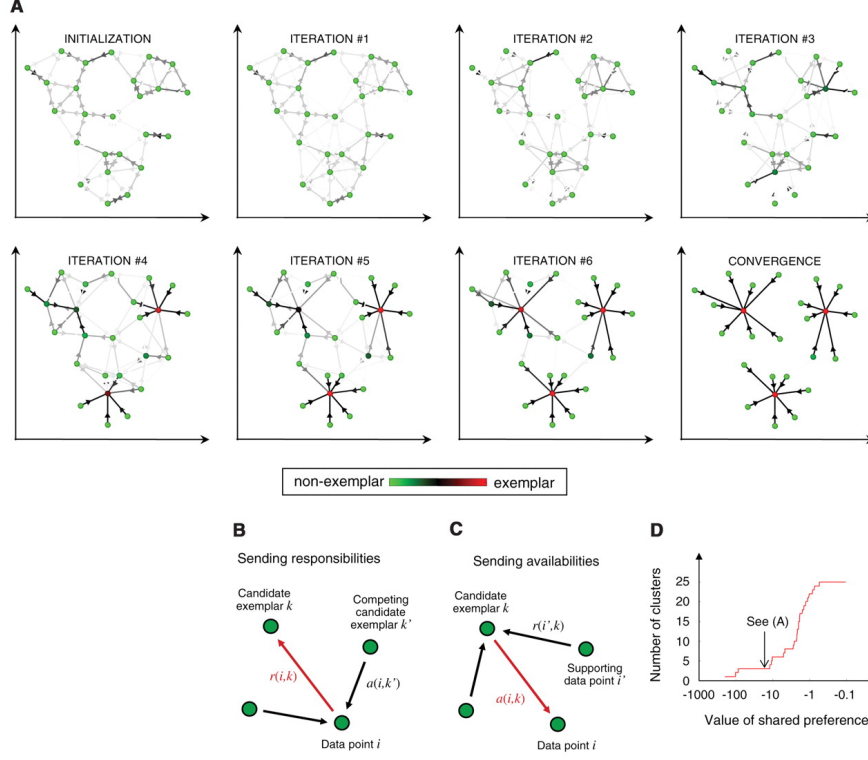


Figure 16: How affinity propagation works

distance) and LCS and L1 combined together through the formula presented in Section 5.

The data fields for each test are defined as follows:

- **Number of scans:** the number of scans that were utilized in the test.
- **Estimated words:** the number of words estimated on the basis of the number of scans(50x).
- **Extracted words:** the number of words actually mined and processed, as well as the corresponding percentage.
- **Number of clusters:** the number of clusters created in process.
- **Features extraction time:** the time, in seconds, required for features extraction from the words.
- **Evaluating distances time:** the time, in seconds, required to generate the similarity matrix with the distances between all the words extracted.
- **Clustering time:** the time, in seconds, required for the creation of clusters.
- **Running time:** the total execution time, in seconds.
- **Correct clusters:** the number of clusters with absolute precision, therefore containing only words equal to each other.

- **Correct elements:** the number of words within the Correct clusters and their corresponding percentage respect to all the elements.
- **Average precision:** the accuracy of the results, the average accuracy of clusters.

$$Precision_{average} = \frac{\sum_i P_i}{N_c}$$

where P_i is the precision of cluster i and N_c is the number of clusters.

- **Precision:** the accuracy of the results, the average accuracy of individual clusters weighted with the number of words.

$$Precision = \frac{\sum_i P_i * n_i}{N_w}$$

where P_i is the precision of cluster i , n_i is the number of words in cluster i and N_w is the number of words.

In the next table we're going to show the main results of the tests.

	Estimated words	Extracted words	Number of clusters	Running time (s)	Precision (%)
16 scans (L1)	800	550	55	15.52	58.00
16 scans (LCS)	800	550	71	44.39	60.00
16 scans (LCS and L1)	800	550	66	47.41	62.91
32 scans (L1)	1600	800	67	38.58	52.87
32 scans (LCS)	1600	800	92	94.54	55.50
32 scans (LCS and L1)	1600	800	86	112.90	56.25
75 scans (L1)	3750	2350	173	221.03	63.65
75 scans (LCS)	3750	2350	189	1169.72	67.49
75 scans (LCS and L1)	3750	2350	199	1203.28	71.16
130 scans (L1)	6500	5050	425	5444.43	71.12
130 scans (LCS)	6500	5050	441	5629.71	74.41
130 scans (LCS and L1)	6500	5050	463	6729.62	77.94
500 scans (L1)	25000	18146	3957	85545.11	77.89
500 scans (LCS)	25000	18146	3316	76324.63	81.22
500 scans (LCS and L1)	25000	18146	3941	138309.85 ⁵	82.14

Table 2: Main results

As we can see in Table 2 the number of extracted words is much less than the estimated number of words. This is mainly due to the fact that not all census tables are filled in full: in some cases there are only a few lines or it has not been filled in the status column. In other cases the error is due to a difficulty in removing the word from the scan due to a wrong interpretation of the rows or the columns.

As we can see in Table 3 and in Figure 17 the accuracy of the cluster grows with the amount of words extracted. This phenomenon is due to the fact that Affinity Propagation works best with a large number of available data: the greater the number of words, the greater the chances of finding words similar between them, and then combine them within a single cluster.

⁵During this test, the machine used was concurrently executing other tasks creating a bottleneck in the allotted memory, in all similar tests the time was in the order of 90k seconds.

	Average precision (%)	Precision (%)
16 scans (L1)	65.06	58.00
16 scans (LCS)	76.56	60.00
16 scans (LCS and L1)	74.37	62.91
32 scans (L1)	59.85	52.87
32 scans (LCS)	72.97	55.50
32 scans (LCS and L1)	70.19	56.25
75 scans (L1)	69.03	63.65
75 scans (LCS)	75.09	67.49
75 scans (LCS and L1)	78.01	71.16
130 scans (L1)	76.66	71.12
130 scans (LCS)	82.11	74.41
130 scans (LCS and L1)	84.87	77.94
500 scans (L1)	90.66	77.89
500 scans (LCS)	91.02	81.22
500 scans (LCS and L1)	92.27	82.14

Table 3: Precision




Figure 17: Clustering precision based on number of scans

Now we're going to focus our attention to time. The running time is divided mainly into three distinct phases: time for the extraction of features, time for the creation of similarity matrix (calculation of distances) and clustering.

	Features extraction time (s)	Evaluating distances time (s)	Clustering time (s)	Total (s)
16 scans (L1)	6.81	5.02	3.69	15.52
16 scans (LCS)	6.84	34.40	3.15	44.39
16 scans (LCS and L1)	6.84	37.32	2.85	47.41
32 scans (L1)	11.54	7.22	19.82	38.58
32 scans (LCS)	11.27	74.72	8.55	94.54
32 scans (LCS and L1)	11.34	89.96	11.60	112.90
75 scans (L1)	33.65	69.00	118.38	221.03
75 scans (LCS)	35.13	1002.63	131.96	1169.72
75 scans (LCS and L1)	34.70	1070.92	97.66	1203.28
130 scans (L1)	70.80	313.70	5059.93	5444.43
130 scans (LCS)	68.17	4353.85	1207.69	5629.71
130 scans (LCS and L1)	73.10	5608.82	1047.70	6729.62
500 scans (L1)	289.33	4482.02	80773.76	85545.11
500 scans (LCS)	323.66	57699.51	18301.46	76324.63
500 scans (LCS and L1)	403.41	61424.57	76481.87 ⁵	138309.85 ⁵

Table 4: Running time

As we can see in Table 4 the computational time is mainly due to the construction phase of similarity matrix, so the calculation of distances. This occurs especially in calculating LCS distance due to the fact that structural strings of words are very long and the cost of the algorithm is $O(nm)$, with n and m the lengths of the two strings.

Structural long strings characterize better the word (and therefore make more accurate clustering), but involve a longer calculation time.




Figure 18: Running time on number of scans

A parameter to evaluate the goodness of the newly formed clusters can be the number of clusters that have a maximum precision, which contain all similar elements between

them and properly classified. This parameter is indicative as it can still exist cases of *false negatives* (words properly classified but labelled incorrectly) that then cut out the entire cluster.

In the following table are collected the number of clusters that respect the above described property and the number of items that they contain.

	Correct clusters	Correct words	Percentage of words (%)	Average words for cluster
16 scans (L1)	12	18	3.27	1.50
16 scans (LCS)	33	44	8.00	1.33
16 scans (LCS and L1)	25	45	8.18	1.80
32 scans (L1)	11	11	1.38	1.00
32 scans (LCS)	39	48	6.00	1.23
32 scans (LCS and L1)	32	43	5.38	1.34
75 scans (L1)	39	172	7.32	4.41
75 scans (LCS)	78	265	11.28	3.40
75 scans (LCS and L1)	77	316	13.45	4.10
130 scans (L1)	138	651	12.89	4.72
130 scans (LCS)	208	713	14.12	3.43
130 scans (LCS and L1)	215	860	17.03	4.00
500 scans (L1)	2854	7354	40.53	2.58
500 scans (LCS)	2902	7423	40.91	2.56
500 scans (LCS and L1)	3020	7616	41.97	2.52

Table 5: Correct clusters

7 Compiling and running notes

This software was developed entirely in C++, and uses an open source library specialized in image processing and analysis, *Leptonica*⁶, version 1.70. Due to technical reasons we couldn't use the newest version 1.71 were even .j2k files are supported, instead we manually converted the files to the already supported .jpg extension and proceeded with them.

Leptonica provides many functions for manipulating images pixel by pixel using a high-level approach. Thanks to this library for example, you can draw up a diagram of projections, crop images, or find the connected components in a portion of the image.

To compile the code, once included the library described above, it is necessary, in the case of version of *GCC/G++* less than 4.7, compile with version 11 of C++ that introduces support for threads.

To run this program we use the following parameters:

- **-d** to specify the images directory, from there the program will automatically load all the files with the specified extension (in our case .jpg) and the corresponding .txt files utilized in the determination of the precision.
- **-t** to specify the number of threads. Default is 2.

⁶Leptonica, a pedagogically-oriented open source site containing software that is broadly useful for image processing and image analysis applications -<http://www.leptonica.org/>

- To determine the distance matrix used (default lcs+l1) one can use:
 - `--lcs` to use only LCS distance for building similarity matrix.
 - `--l1` to use only Euclidean distance for building similarity matrix.

8 Conclusion

In this project we wrote an application that after accepting as input U.S. Census documents returns clusters containing words that corresponds to the same state of birth. Starting from a pre-existing work that already located and segmented the samples corresponding to the citizens' state, we added a more thorough search for the targeted area, and with the intent of improving the application by providing an alternate method of calculating the distance matrix we developed a different kind of features that we feel are more closely associated to the way a word is written through different strokes. Through the use of these *primitive* features and the already defined dimensional features we are able to generate distance matrices based only on one of the features' kind or on both after a reasoned combination. We utilize then the Affinity Propagation algorithm to obtain the desired clusters.

Based on the results obtained we can note that the accuracy of the cluster grows with the amount of words extracted. This phenomenon is due to the fact that Affinity Propagation works best with a large number of available data: the greater the number of words, the greater the chances of finding words similar between them, and then combine them within a single cluster. At the same time, with the increase of the processed words there's also an increase in the rate of correct clusters.

We note that the time needed to complete the calculations increases quadratically with the number of elements. This, as shown in Table 4, mainly due to the time needed by LCS to calculate the distance between each pair of words. Using other calculation methods one might reduce the necessary computing time.

From the tests the new features clearly provide an improvement on the older dimensional features, this improvement can further be expanded through the use of the mixed distance matrix. The improvement is sharp when low numbers of scans are considered but decreases at higher numbers where the precision of the 3 methods tend to converge.

The application overall succeeds in our intent although with not incredible improvements on the preceding work, possible avenues of optimization are in the definition of the strings associated to the features, taking in consideration that an increase in the string length is directly proportional to an increase in LCS distance's calculation time.

Other possible and obvious improvement, although that would require an overall rewriting of the project from it's foundation, would be the association of each segmented word with its coordinates in the general census image: this cannot be done at the moment because the coordinates of a segment are ignored since the preexisting inherited segmentation steps that simply cuts the words without any reference to their place in the general image. This improvement would confer an effective practical utility to the application; at the same time it would permit a betterment of the debug and testing phase through the possible use of the same coordinates in the search of the word corresponding to the segment, rather than having to rely on each segment's row number.